

Jefrey Kerley
Jakcqc
2/7/2022

Problem 1:

A) The difference between test MSE and Training MSE

is that the test MSE is the MSE as it relates to more general data.

Meaning, that you would use the training MSE to gauge how well

the line of best fit is, only when dealing with the sample data.

So the difference between the two MSEs, is that one is relevant to the

sample data, and one is relevant to new general data, that was not in

the training set.

B) The two curves show a different trend because as the complexity for

training set increases, the line over fits and is able to "correctly"

fit all training data. Yet, when this is applied to the testing data,

we see that the error sky rockets as the over fitted line does poorly

on data that does not match the training data.

C) The large test MSE for the linear regression fit, is that the line

was too generalized, and because the "true" line of best fit was a

cubic line. As a result, more complexity needed to be added to prevent

a large testing MSE for this line.

D) The spline with high degrees of freedom has a high test MSE, because it

is not generalized at all. This is a classic over fitted line, that was

trained "too much", on the training data. instead of learning, this line

memorized the data.

E) The training MSE that are lower than the irreducible error, as this error

acts as the minimum lower bound for the test MSE. Which is why we see the

test MSE never drop below this. Over fitted, and "perfect" lines

for some data, it is possible for the test MSE to drop practically to 0.

Furthermore, this error simply acts as a representation of the real world

truth that is unknown, ie- known knowns, unknown knowns, etc...

##

Problem 2

a)

```
collegeData = read.csv('./College.csv', stringsAsFactors = TRUE)
```

```
View(collegeData)
```

b)

```
rownames(collegeData) = collegeData[, 1]
```

```
View(collegeData)
```

```
collegeData = collegeData[, -1]
```

```
View(collegeData)
```

c)

```
summary(collegeData)
```

```
pairs(collegeData[,1:10])
```

```

plot(collegeData$Private, collegeData$Outstate)
Elite <- rep ("No", nrow (collegeData))
Elite[collegeData$Top10perc > 50] <- " Yes "
Elite <- as.factor(Elite)
collegeData <- data.frame (collegeData , Elite)
summary(Elite)
plot(Elite, collegeData$Outstate)
par(mfrow = c(2,2))
hist(collegeData$Top10perc)
hist(collegeData$Top25perc)
hist(collegeData$Enroll)
hist(collegeData$Apps)
##
## Problem 3
##
?data.frame
?runif
set.seed(80)
train.X <- data.frame(x1 = runif(n=100,-1,1), x2 = runif(n=100,-1,1))
# Randomly select 100 locations of (x1, x2)
prob <- ifelse( (-0.5 < train.X$x1) & (train.X$x1 < 0.5)
               & (-0.5 < train.X$x2) & (train.X$x2 < 0), 0.9, 0.1)
# If (x1, x2) is inside the rectangle, prob is 0.9, otherwise, prob is 0.1
train.Y <- as.factor(runif(n=100) < prob)
# Simulate class labels according to prob
colors <- c("skyblue","red")
plot(train.X$x1, train.X$x2, pch=20, col=colors[factor(train.Y)])
segments(-0.5, -0.5, 0.5, -0.5, col="orange", lwd=2)
segments(-0.5, -0.5, -0.5, 0, col="orange", lwd=2)
segments(-0.5, 0, 0.5, 0, col="orange", lwd=2)
segments(0.5, -0.5, 0.5, 0, col="orange", lwd=2)
## a) The Bayes error rate would be .10, as the irreducible error == bayes
## error, thus are prediction rate would be 90% given a large set of data.
## To further explain, Bayes error is the minimum error you can achieve.
##
## b) Below we can see that the distribution is a reasonable result.
set.seed(100)
train.X <- data.frame(x1 = sample(0:1, 100, replace = T))
train.Y <- as.factor(train.X$x1 > 0)
colors <- c("black","blue")
plot(train.Y, col=colors[factor(train.Y)], ylim = range(0,70))
##
## c) Below we can see that when k = 1, it does not satisfy the predictive
## outcome. The classification "areas", are not relative to the right

```

```

## classification predictive classes.
xGrid.1d <- seq(-1,1,0.02)
nx <- length(xGrid.1d)
xGrid.2d <- data.frame(x1=rep(xGrid.1d, each=nx), x2=rep(xGrid.1d,times=nx))
library(class)
set.seed(1)
runBlockKNN = function(knnSize)
{
  knn.pred <- knn(train=train.X, test=xGrid.2d, cl=train.Y, k=knnSize)
  plot(xGrid.2d$x1, xGrid.2d$x2, pch=3, col=colors[factor(knn.pred)])
  segments(-0.5, -0.5, 0.5, -0.5, col="orange", lwd=2)
  segments(-0.5, -0.5, -0.5, 0, col="orange", lwd=2)
  segments(-0.5, 0, 0.5, 0, col="orange", lwd=2)
  segments(0.5, -0.5, 0.5, 0, col="orange", lwd=2)
}
runBlockKNN(1)

```

d) With increasing values of k, to an upper bound, we see the classification
 ## becoming more true to the right predicitive classification.

```

runBlockKNN(2)
runBlockKNN(3)
runBlockKNN(4)
runBlockKNN(5)
runBlockKNN(6)
runBlockKNN(7)
runBlockKNN(8)
runBlockKNN(12)
runBlockKNN(15)

```

e) The best k value here seems to be around k = 10, up to maybe around
 ## k = 15. This is where the classification starts to be more accurate
 ## of the expective predictive outcome.

```

##
##
##
##
##
##
##
##
##
##
##

```